

CARECHRONICLE

Professional Product Requirements Document

Patient-Owned Medical Record & Care Organization Platform

Turning scattered medical documents into one organized, longitudinal, actionable medical story.

Field	Value
Document Version	2.0
Product Name	CareChronicle
Document Type	Professional Product Requirements Document (PRD)
Primary Launch Focus	Cancer care / complex longitudinal care
Platform	Mobile-first application + responsive web experience
Status	Product definition / pre-development

Core Product Principle: AI is the engine; the persistent medical-record workflow is the product.

1. Executive Summary & Product Definition

CareChronicle is a patient-owned digital medical record and care-organization platform for people managing cancer and other complex, longitudinal healthcare journeys. It is designed around a practical reality: medical information accumulates across hospitals, laboratories, pharmacies, imaging centres, doctors, insurers, email, phones and physical paperwork.

The core problem is not merely document storage. Patients and caregivers need to know what happened, when it happened, which documents belong together, what changed, where the original source is, and what needs to be shared or submitted next.

CareChronicle therefore creates a persistent organizational layer around the patient's records. It collects documents, preserves originals, classifies and extracts metadata, links records to appointments and care episodes, builds a longitudinal timeline, enables retrieval and produces purpose-specific packages for appointments, second opinions and insurance administration.

Product promise: a patient should not have to remember which folder contains the prescription from a particular visit or which receipt belongs to a hospitalization. CareChronicle should make those relationships explicit.

The supplied research identifies information fragmentation, longitudinal context loss, medication discrepancies, care-transition complexity, health-literacy challenges, symptom communication and navigation burden as recurring issues in cancer care. It also identifies existing adjacent products, meaning differentiation must come from workflow depth, reliability and trust—not from generic AI summarization alone.

2. Problem Statement & Opportunity

A complex medical journey can generate consultation notes, pathology, imaging, blood tests, prescriptions, treatment records, discharge summaries, referrals, hospital invoices, pharmacy receipts, payment records and insurance documents. These records often arrive as PDFs, scans, photographs, portal downloads, emails or paper.

When these artifacts remain disconnected, users repeatedly perform manual administrative work: searching, renaming, sorting, reconstructing chronology, asking relatives for documents, preparing doctor packets and assembling insurance paperwork.

Opportunity: build the missing patient-owned information layer that connects the records without pretending to become the clinical decision-maker.

2.1 Core User Problems

ID	Problem	Description
P-01	Scattered records	Documents live across multiple hospitals, labs, phones, email, portals and paper files.
P-02	Poor chronology	Users struggle to reconstruct the care journey from disconnected files.
P-03	Weak relationships	A prescription, report and bill from one visit may be stored independently.
P-04	Repeated searching	Old records must be located again for visits, referrals, second opinions and claims.
P-05	Longitudinal context loss	Historical information is difficult to compare with new records.
P-06	Caregiver burden	Family members often become the informal medical-record managers.
P-07	Second-opinion friction	Relevant records must be manually selected and arranged.
P-08	Insurance administration	Financial and medical documents must be found and packaged together.
P-09	AI trust risk	OCR/extraction can be wrong; users need verification and source traceability.

3. Product Vision, Mission, Goals & Non-Goals

Vision: Every patient should be able to carry a clear, organized and usable record of their healthcare journey wherever they receive care.

Mission: Make complex medical information easier to collect, organize, retrieve, share and use—without replacing clinical judgment.

3.1 Product Goals

- Create a patient-owned source of truth for uploaded medical documents.
- Automatically organize records into appointment and care-episode structures.
- Preserve original documents while generating searchable structured metadata.
- Reduce time and cognitive effort required to prepare for appointments.
- Make longitudinal context easier to retrieve and review.
- Generate organized documentation packages for second opinions and insurance workflows.
- Enable secure caregiver collaboration with granular permissions.
- Build trustworthy AI workflows with visible provenance and human confirmation.

3.2 Non-Goals

- Diagnosing cancer or any other disease.
- Recommending, changing or discontinuing treatment or dosage.
- Predicting survival or treatment outcomes.
- Replacing doctors, pharmacists, nurses or emergency services.
- Guaranteeing insurance eligibility or claim approval.
- Being an emergency response or clinical triage system.
- Integrating deeply with every hospital/EHR in the MVP.
- Trying to support every disease equally at initial launch.

4. Target Users, Personas & Jobs-to-be-Done

Persona	Description	Primary Needs
Primary Patient	Person receiving complex care.	Own records, find documents, prepare for visits, share history, manage documentation.
Caregiver	Family member supporting the patient.	Upload, organize, retrieve, prepare summaries and packages with permission.
Clinician Recipient	Doctor receiving patient-provided records.	Receive concise, relevant context and source documents.
Administrative User	Person preparing financial/insurance documentation.	Collect and package related records efficiently.

4.1 Jobs-to-be-Done

- When I receive a new medical document, I want it filed correctly without manually creating folders.
- When I have an appointment, I want everything relevant to that visit immediately available.
- When I meet a new doctor, I want to share an organized history rather than dozens of random files.
- When I receive a new report, I want to find the relevant previous record quickly.

- When I prepare an insurance claim, I want related medical and financial records identified and organized.
- When a family member helps me, I want them to have the right access without losing control.

4.2 Representative User Stories

ID	Actor	Action	Expected Outcome
US-01	Patient	Bulk upload documents	System classifies documents and proposes filing.
US-02	Patient	Open an appointment	See reports, prescriptions, notes and financial records linked to it.
US-03	Caregiver	Upload a receipt	Receipt becomes searchable and linkable after verification.
US-04	Patient	Prepare for visit	Generate a concise visit packet and recent-change view.
US-05	Patient	Prepare second opinion	Generate chronological, organized source-document package.
US-06	Patient	Prepare insurance file	Find and package relevant uploaded documentation by episode/date range.

5. Information Architecture & Core Data Model

The central design decision is to treat the **appointment/care episode** as the organizing unit while retaining a global Medical Vault. A document may belong to an appointment, hospitalization, treatment cycle, diagnostic episode or another care event.

5.1 Primary Navigation

- **Home:** recent activity, upcoming appointments, quick actions and unfiled/verification tasks.
- **Medical Vault:** all original and generated records with filters and previews.
- **Timeline:** chronological care history.
- **Appointments / Episodes:** structured dossiers of related records.
- **Search:** metadata, full-text and semantic retrieval.
- **Packages:** doctor, second-opinion and insurance exports.
- **People & Access:** caregiver and sharing permissions.
- **Settings:** privacy, security, export, deletion and account controls.

5.2 Appointment Folder Example

04 September 2026 — Oncology Consultation

- Consultation note / doctor note
- CBC / LFT / pathology / other laboratory reports
- Imaging report
- Prescription
- Treatment instructions
- Hospital / diagnostic bills
- Pharmacy receipts
- Follow-up instructions
- Patient questions / notes

5.3 Core Entities

Entity	Definition	Purpose
Patient	Identity and profile metadata.	Owns cases and permissions.
Case	Longitudinal medical journey grouping.	Top-level container.
Care Episode	Hospitalization, consultation, diagnostic event, treatment cycle, etc.	Groups related activity.
Appointment	Specific visit or consultation.	Primary filing context.
Document	Original uploaded record.	Immutable source artifact.
Extraction	AI-derived structured information.	Traceable to source.
Document Link	Relationship between a document and case/episode/appointment.	Connects context.
Package	Generated export bundle.	Doctor, second opinion or insurance.
Access Grant	Permission assigned to another person.	Controls collaboration.
Audit Event	Recorded sensitive action.	Accountability and security.

6. Functional Requirements

ID	Priority	Requirement	Acceptance Summary
FR-01	M	Account & Case	Create account, patient profile and medical cases.
FR-02	M	Document Upload	Single and bulk PDF/image upload; camera capture.
FR-03	M	Original Preservation	Original file retained without destructive modification.
FR-04	M	OCR/Text	Extract searchable text from scans and digital PDFs.
FR-05	M	Classification	Identify common document categories.
FR-06	M	Metadata	Extract dates, provider, facility and administrative metadata.
FR-07	M	Verification	Show uncertainty and allow correction before filing.
FR-08	M	Appointment Filing	Suggest appointment/episode association and allow correction.
FR-09	M	Medical Vault	Browse, filter, preview, download and manage records.
FR-10	M	Timeline	Generate chronological events with source links.
FR-11	M	Search	Search by text, date, type, provider, medication and episode.
FR-12	M	Appointment View	Display all linked records in one dossier.
FR-13	M	Visit Summary	Generate concise, source-grounded appointment summary.
FR-14	M	Export / Share	Generate controlled PDF/ZIP packages and secure links.
FR-15	M	Audit Trail	Record important uploads, edits, shares and exports.
FR-16	S	What Changed	Compare selected records and surface record-level differences.
FR-17	S	Second Opinion	Generate organized medical-history package.
FR-18	S	Insurance Package	Collect relevant financial/medical documents by episode/date range.

ID	Priority	Requirement	Acceptance Summary
FR-19	S	Caregiver Roles	Invite caregivers with granular permissions.
FR-20	C	Advanced Integrations	Future hospital/EHR/insurer integrations.

6.1 Document Ingestion Workflow

1. User selects one or multiple files or captures a document.
2. System securely stores the original.
3. OCR/text extraction runs where required.
4. Classifier proposes document type.
5. Metadata extraction proposes dates, provider and facility.
6. System proposes an appointment/care-episode association.
7. Confidence indicators identify uncertain fields.
8. User confirms or corrects uncertain information.
9. Document becomes searchable and visible in the selected case.
10. System records AI provenance and processing version.

6.2 Edge Cases

- Document date differs from upload date.
- Lab collection date differs from report publication date.
- One PDF contains multiple reports.
- One appointment includes multiple providers/facilities.
- Duplicate documents are uploaded.
- Document is handwritten or low quality.
- Patient name is abbreviated or inconsistent.
- Document is bilingual or multilingual.
- User is unsure which appointment a document belongs to.
- Classification/extraction fails.

7. AI & Document Intelligence

AI exists to reduce manual organization work—not to make clinical decisions. Material AI-generated information must remain traceable to the source document.

ID	Capability	Requirement
AI-01	OCR	Convert scanned pages/images into searchable text.
AI-02	Classification	Identify document category.
AI-03	Metadata Extraction	Extract date, provider, facility and document title.
AI-04	Entity Extraction	Identify information explicitly present in source documents.
AI-05	Relationship Inference	Propose links between documents and care events.
AI-06	Summarization	Create concise, source-grounded administrative/visit summaries.
AI-07	Comparison	Highlight differences between selected records.

ID	Capability	Requirement
AI-08	Semantic Search	Retrieve relevant records using natural-language queries.

7.1 AI Trust Requirements

- Never fabricate a date, medication, test result, document or event.
- If evidence is unavailable, explicitly state that it is unavailable.
- Preserve source-document links for extracted facts.
- Show confidence/verification status for uncertain extraction.
- Allow user corrections.
- Never overwrite the original record.
- Label generated summaries as AI-generated.
- Use deterministic rules where they are safer than generative models.

7.2 Clinical Safety Boundary

Allowed: organization, extraction, summarization, retrieval, comparison, document packaging and administrative assistance.

Not allowed: diagnosis, treatment recommendation, dosage adjustment, medication discontinuation instructions, prognosis, emergency triage or claims that a medical decision is safe/unsafe.

8. Core User Workflows

8.1 Appointment Preparation

1. Open an appointment or create one.
2. Review linked documents and newly added records.
3. Review suggested related records.
4. Generate 'Since Your Last Visit' summary.
5. Review and edit summary if needed.
6. Open source documents for verification.
7. Export/share a concise visit package.

8.2 What Changed

Compare two selected reports or periods and surface record-level changes such as changed values, changed statements, newly present information or prescription changes. The UI must distinguish source facts from interpretation and must not convert a change into a diagnosis or treatment recommendation.

8.3 Second Opinion Package

- Select case and relevant date range.
- Recommend relevant categories: pathology, imaging, labs, treatment history, prescriptions, discharge summaries and notes.
- Preview and remove irrelevant documents.
- Generate chronological index.
- Export original source documents.
- Create controlled share link with scope and expiry.

8.4 Insurance Documentation Package

- Select treatment/hospitalization/episode and date range.
- Find related bills, receipts, prescriptions, investigations, admission/discharge records and payment documents.
- Show categories found and categories for which no matching uploaded record was found.
- Allow manual additions/corrections.
- Generate organized index plus original documents.
- Do not guarantee insurer acceptance or infer policy-specific requirements without verified policy information.

9. UX / UI Requirements

CareChronicle should look and feel like a calm, trustworthy medical filing system—not a flashy chatbot. The design language should emphasize hierarchy, source traceability, clear status and low cognitive load.

ID	Area	Requirement
UX-01	Home	Show cases, recent uploads, appointments and verification tasks.
UX-02	Upload	Support file picker, camera scan and bulk import.
UX-03	Verification	Show uncertain fields beside source preview.
UX-04	Appointment	Dossier-style view with clear categories.
UX-05	Timeline	Chronological events with expandable source links.
UX-06	Search	Natural language + structured filters.
UX-07	Sharing	Recipient, scope, permissions and expiry are explicit.
UX-08	Trust	AI-generated information visibly labeled with 'view source'.

9.1 Accessibility

- Large touch targets and readable typography.
- High contrast for essential information.
- Do not rely on color alone for status.
- Plain-language labels and progressive disclosure.
- Screen-reader-compatible semantics.
- Design for older adults and caregivers.

9.2 Localization

Architecture should support English first and future Indian-language interfaces. Dates, currency, provider names and document formats should be locale-aware. Regional-language OCR should be validated before broad support claims.

10. Security, Privacy, Compliance & Trust

Medical records are highly sensitive. Security and privacy are core product requirements.

ID	Area	Requirement
SEC-01	Encryption	Encrypt data in transit and at rest; protect keys separately.
SEC-02	Authentication	Strong authentication with optional MFA/biometric unlock.

ID	Area	Requirement
SEC-03	Authorization	Least-privilege access at patient/case/document levels.
SEC-04	Sharing	Explicit recipient, scope and expiry controls.
SEC-05	Audit	Log important access, edit, export and share events.
SEC-06	Data Minimization	Collect only data required for functionality.
SEC-07	Deletion/Export	User-controlled export/deletion subject to applicable retention rules.
SEC-08	AI Processing	Define model-provider data use, retention and processing location.
SEC-09	Backups	Encrypted backups with restore testing.
SEC-10	Incident Response	Document detection, containment, notification and recovery.

10.1 Compliance Direction

The product should be designed with applicable Indian privacy/data-protection obligations in mind and reviewed by qualified legal/compliance professionals before production launch. Expansion to other jurisdictions requires a separate assessment.

Trust requirement: users must always be able to distinguish original medical records from information generated or extracted by CareChronicle.

11. Technical Architecture

Layer	Technology Direction	Responsibilities
Client	Mobile-first + responsive web	Authentication, upload, vault, timeline, search, sharing, packages.
API	Backend services	Identity, cases, documents, appointments, permissions, exports, audit.
Document Service	Object storage + metadata DB	Original files, versions, checksums, lifecycle.
Processing Pipeline	Async queue/workers	Scan → OCR → classify → extract → index → verify.
AI Layer	Model abstraction	LLM/NLP services with provider isolation and evaluation.
Search	Hybrid search	Metadata + full text + semantic retrieval.
Database	Relational DB	Cases, episodes, appointments, documents, links, permissions.
Export	Package generator	PDF index, ZIP bundle, secure links and expiry.
Observability	Logs/metrics/tracing	Reliability, security events and product analytics.

11.1 Processing Pipeline

Upload → security scan → normalization → OCR → classification → metadata/entity extraction → appointment/episode matching → confidence scoring → human verification → indexing → timeline update → package/search availability.

11.2 Reliability Strategy

- Every asynchronous job is retryable and idempotent.
- Failed AI/OCR jobs never prevent retention of the original.
- Document access remains usable even if AI processing is unavailable.

- Processing status is visible to the user.
- Originals remain immutable; derived information can be reprocessed.

12. Data Model & API Requirements

12.1 Core Data Objects

Object	Core Fields / Purpose
User	Authentication, preferences and security settings.
PatientProfile	Identity and administrative metadata.
Case	Longitudinal case/condition grouping.
CareEpisode	Hospitalization, treatment cycle, diagnostic or consultation event.
Appointment	Visit details and linked episode.
Document	Storage key, checksum, MIME type, upload/source metadata.
DocumentMetadata	Type, dates, provider, facility, confidence, verification state.
Extraction	Field, value, source location, model version, confidence.
DocumentLink	Relationship between document and case/episode/appointment.
Package	Type, selected documents, creation time, expiry.
AccessGrant	Recipient, scope, permissions, expiry, status.
AuditEvent	Actor, action, object, timestamp and metadata.

12.2 API Surface

Domain	Endpoint Pattern	Purpose
Auth	POST /auth/*	Sign-up, login and session management.
Patients	/patients/*	Profile and case ownership.
Documents	/documents/*	Upload, metadata, preview, download and delete.
Processing	/documents/{id}/processing	Processing status and reprocessing.
Appointments	/appointments/*	Create, link and browse appointments.
Timeline	/cases/{id}/timeline	Chronological event retrieval.
Search	/search	Hybrid search and filtering.
Packages	/packages/*	Create, preview, export and share.
Access	/access/*	Caregiver/recipient permissions.
Audit	/audit/*	Authorized audit retrieval.

13. Non-Functional Requirements

Area	Requirement
Performance	Common navigation should feel responsive; heavy processing is asynchronous.

Area	Requirement
Availability	Define production SLOs; document access should remain resilient during AI outages.
Scalability	Support large libraries and bulk uploads without redesign.
Reliability	No loss of originals; checksums and backup validation.
Security	Security testing, dependency scanning and least-privilege access.
Privacy	Clear consent, data-use disclosure, export and deletion controls.
Observability	Centralized errors, processing metrics and audit logs.
Accessibility	WCAG-informed web design and mobile accessibility guidelines.
Localization	Locale-aware dates/currency and future multilingual readiness.
Maintainability	Versioned AI prompts/models, migrations and modular services.

13.1 Example Engineering Targets

Targets should be finalized during technical discovery. Candidate targets include $\geq 99.5\%$ successful uploads excluding user/network failures, zero intentional loss of originals, searchable availability after successful processing, defined API error-rate SLOs, and source traceability for all user-visible generated summaries.

14. Analytics, Metrics & Success Criteria

CareChronicle should measure whether it reduces organizational work rather than optimizing for superficial daily engagement.

Metric	Definition
Activation	% of users who upload and successfully organize their first meaningful document set.
Organization Success	% of uploaded documents correctly filed after verification.
Verification Burden	Average manual corrections per processed document.
Retrieval Efficiency	Median time required to find a requested document.
Appointment Readiness	% of appointments with an organized pre-visit package.
Package Completion	% of packages completed without external manual organization.
Sharing Success	% of generated packages successfully accessed by intended recipient.
Trust	User-rated confidence in organization and source traceability.
Care-Event Retention	% of users returning when a new care event occurs.
Support Rate	Support tickets per active patient/case.

14.1 Key Events

- case_created
- document_uploaded
- document_classified
- verification_requested
- verification_completed
- appointment_created

- document_linked
- search_performed
- document_opened_from_search
- summary_generated
- package_created
- package_shared
- caregiver_invited
- document_exported
- document_deleted

15. QA, Validation & Acceptance

Document intelligence is the hardest component. Testing must combine conventional software QA with a dedicated AI/document evaluation program.

Test Domain	Coverage
Functional QA	Core flows, permissions, upload, browse, search, export and sharing.
Document QA	PDFs, scans, photos, multi-page reports, duplicates and corrupt files.
OCR QA	Printed text, low-resolution scans, skew, tables and representative languages.
Extraction QA	Dates, providers, facilities, document types, medications and administrative fields.
Relationship QA	Correct and incorrect appointment/episode associations.
AI Safety QA	Hallucination, unsupported inference, clinical-advice leakage and provenance.
Security QA	Authentication, authorization, IDOR, session management, encryption and sharing.
Performance QA	Bulk upload, concurrent processing, large libraries and export generation.
Usability QA	Older adults, caregivers and first-time users.

15.1 MVP Acceptance Criteria

- User can create a case and upload individual or bulk documents.
- Original files remain downloadable.
- Common document categories are classified at a measured target accuracy.
- Extracted metadata is visible and correctable.
- Appointment/episode suggestions can be corrected.
- Timeline is generated with source links.
- Search works across metadata and text.
- Appointment view shows linked records.
- Visit summary is source-grounded.
- Selected documents can be exported with an index.
- Caregiver access is permission-controlled and auditable.
- AI-generated information is clearly distinguished from source records.

16. MVP Scope, Roadmap & Delivery

16.1 MVP — Must Ship

- Secure account and case creation.
- Medical Vault.
- Bulk PDF/image upload and camera capture.
- OCR and basic document classification.
- Metadata extraction with human verification.
- Appointment/care-episode organization.
- Longitudinal timeline.
- Basic smart search.
- Appointment view.
- Source-grounded one-page visit summary.
- Basic export package.
- Caregiver access foundation.
- Security and audit foundation.

16.2 Phase 2

- What Changed comparison
- Second-opinion package builder
- Insurance package builder
- Advanced semantic search
- Duplicate detection
- Improved multilingual OCR
- Expanded caregiver workflows

16.3 Phase 3

- Provider/hospital integrations
- Advanced longitudinal analytics
- Medication-history reconciliation
- Structured lab trends
- Interoperability
- Enterprise/provider-facing workflows after validation

16.4 Delivery Sequence

Phase	Theme	Deliverables
1	Foundation	Auth, cases, vault, storage, upload, security baseline.
2	Document Intelligence	OCR, classification, extraction, verification.
3	Organization	Appointments, episodes, links, timeline.
4	Retrieval	Search, filters, previews, relationships.
5	Outputs	Visit summary, export, sharing.
6	Hardening	QA, security, accessibility, analytics, pilot.

17. Risk Register

ID	Risk	Severity	Mitigation
R-01	OCR/extraction errors	High	Confidence scoring, verification, source traceability and evaluation datasets.
R-02	Wrong appointment association	High	Suggest rather than silently assign; multiple date signals; correction UI.
R-03	User distrust	High	Original-first design, provenance and clear AI labeling.
R-04	Privacy/security incident	Critical	Threat modeling, encryption, least privilege, audits and security testing.
R-05	Overly broad MVP	High	Launch around one high-value population and one core workflow.
R-06	Low ongoing engagement	Medium	Optimize around recurring care events and high-value moments.
R-07	Competitive overlap	High	Validate exact feature overlap and differentiate through workflow/trust.
R-08	Regulatory ambiguity	Medium	Legal/compliance review before launch and expansion.
R-09	AI vendor dependency	Medium	Model abstraction, fallback strategy and vendor review.
R-10	Insurance complexity	Medium	Documentation organizer positioning; no approval promises.

18. Competitive Positioning & Product Moat

The supplied research identifies adjacent categories including symptom trackers, medication apps, patient portals, structured treatment summaries and emerging longitudinal-record products. CareChronicle must therefore avoid broad claims such as 'no app manages medical records.'

Positioning: CareChronicle is a patient-owned medical record organization platform for complex care that connects documents to the events that created them, preserves source records, and turns the organized record into ready-to-use appointment, second-opinion and insurance packages.

18.1 Differentiation

- **Appointment-aware filing:** documents are organized around care events, not only generic folders.
- **Longitudinal continuity:** the same patient-owned record persists across providers.
- **Source-grounded intelligence:** generated information is linked to originals.
- **Output workflows:** the product produces packages, not merely summaries.
- **Human-in-the-loop trust:** uncertain extraction is surfaced.
- **Caregiver collaboration:** designed around real family coordination.

Moat hypothesis: the strongest long-term advantage is reliable ingestion of messy real-world medical documents, a high-quality longitudinal data model, trustworthy source relationships and repeated workflow utility—not access to a generic language model.

19. Business Model & Go-to-Market Hypotheses

Commercial choices require validation. The following is a working hypothesis.

Tier	Potential Value	Purpose
Free / Starter	Limited storage and core organization.	Acquisition and workflow validation.
Premium Patient	Higher storage, advanced search, packages and caregiver features.	Monetize high-value organization.
Family / Caregiver	Multiple collaborators and expanded management.	Serve complex family care.
Institutional / B2B2C	Potential provider/insurer partnerships.	Future distribution after validation.

19.1 Initial Go-to-Market Wedge

Start with patients and caregivers managing complex cancer journeys, where document volume, recurring appointments and multi-provider care make the organizational problem especially visible. Expand only after validating the core workflow.

19.2 Distribution Hypotheses

- Cancer support communities and caregiver networks.
- Patient education and hospital-adjacent organizations.
- Referral through care coordinators where appropriate.
- Insurance/documentation assistance workflows.
- Content showing real before/after organization problems.

20. Open Questions & Validation Plan

ID	Question	Validation Method
Q-01	How much manual confirmation is acceptable?	Usability testing.
Q-02	Which document categories are essential at launch?	Pilot dataset analysis.
Q-03	What Indian-language OCR is needed?	User/market research.
Q-04	Should cases be disease-, episode- or patient-oriented?	UX validation.
Q-05	What insurance package is actually useful?	User/insurer interviews.
Q-06	What retention/deletion model is appropriate?	Legal/security review.
Q-07	Which AI vendors satisfy privacy/reliability requirements?	Technical/security assessment.
Q-08	Which competitors overlap with the MVP?	Current market research.
Q-09	What pricing model is viable?	Pricing research.

21. Research & Design Basis

This PRD uses the research and product discussions supplied in the project materials. Source-derived themes include information fragmentation across cancer care, longitudinal context loss, medication discrepancies during transitions, patient-navigation needs, health-literacy challenges, symptom communication, caregiver burden and the value of structured treatment summaries.

The supplied research ranked a longitudinal treatment timeline and one-page summary as a strong opportunity while also identifying adjacent products. It recommended a realistic MVP built around document ingestion, organization, timeline and summary workflows rather than an undifferentiated AI report explainer.

Important: competitive claims, legal/compliance requirements, pricing, clinical boundaries and implementation decisions should be revalidated before production launch.

22. Final Product Definition

CareChronicle is a **patient-owned medical record and care organization platform** that turns fragmented medical documents into a structured, longitudinal record.

Its primary organizing unit is the **care event**—especially the appointment or episode—so reports, prescriptions, notes, bills and receipts can be understood in context. Its AI layer automates document ingestion, classification, extraction, relationship suggestions, search and source-grounded summarization. Its workflow layer converts the organized record into practical outputs: appointment preparation, 'what changed' views, second-opinion packages and insurance documentation packages.

The product succeeds when a patient can answer, without searching through folders or messaging multiple family members: **What happened? Where is the document? What belongs to this visit? What do I need to share or submit now?**

CareChronicle should therefore be built as a **trusted information infrastructure for the patient's care journey**, not as another generic AI medical chatbot.

CARECHRONICLE

Organize the records. Preserve the story. Simplify the journey.